

732A75 Advanced Data Mining laboratory 1 report

Yuki Washio and Nicolas Taba

March 4, 2021

1 Introduction

The aim of this laboratory exercise is to gain familiarity with the data mining toolkit Weka by using some of its clustering algorithms and analyzing their output.

In this laboratory, we use the HARTIGAN (file06) dataset. This dataset features 27 kinds of food and informs us of the energy, protein, fat, calcium and iron content per 3 ounce of portion.

2 Clustering using the Kmeans algorithm

In this section of the laboratory, we use the Kmeans algorithm to cluster the data. We choose to ignore the "name" label because the algorithm does not interpret names as words but as strings. The Kmeans algorithm would try to measure distances between strings which is not a good metric. However, we notice that all of the food are meat and seafood based.

We experiment with the clustering algorithm by trying 2 different number of cluster (2 and 5) starting with the same initial cluster center (seed 10).

2.1 2 clusters with seed 10

The clustering algorithm separates into two clusters with 9 and 18 elements respectively in each Fig. 1. We observe a linear relationship between fat as a function of energy Fig. 2a. We seem to be able to separate the clusters into High energy, high fat the second cluster having low values for those labels. The clusters seem to be made up of similar elements that discriminate mostly along the energy values as we can see in the graphs that depict Protein, calcium and iron as a function of energy respectively.

```

Filename:      kmeans_seed10_2c1.model
Scheme:weka.clusterers.SimpleKMeans -N 2 -A "weka.core.EuclideanDistance -R first-last" -I 500 -S 10
Relation:food
Attributes:6
    Energy
    Protein
    Fat
    Calcium
    Iron
Ignored:
    Name

=== Clustering model ===

KMeans
=====

Number of iterations: 2
Within cluster sum of squared errors: 5.069321339929419
Missing values globally replaced with mean/mode

Cluster centroids:

```

Attribute	Full Data (27)	Cluster#	
		0 (9)	1 (18)
Energy	207.4074	331.1111	145.5556
Protein	19	19	19
Fat	13.4815	27.5556	6.4444
Calcium	43.963	8.7778	61.5556
Iron	2.3815	2.4667	2.3389

Figure 1: Kmeans algorithm output from Weka for 2 clusters, using seed 10.

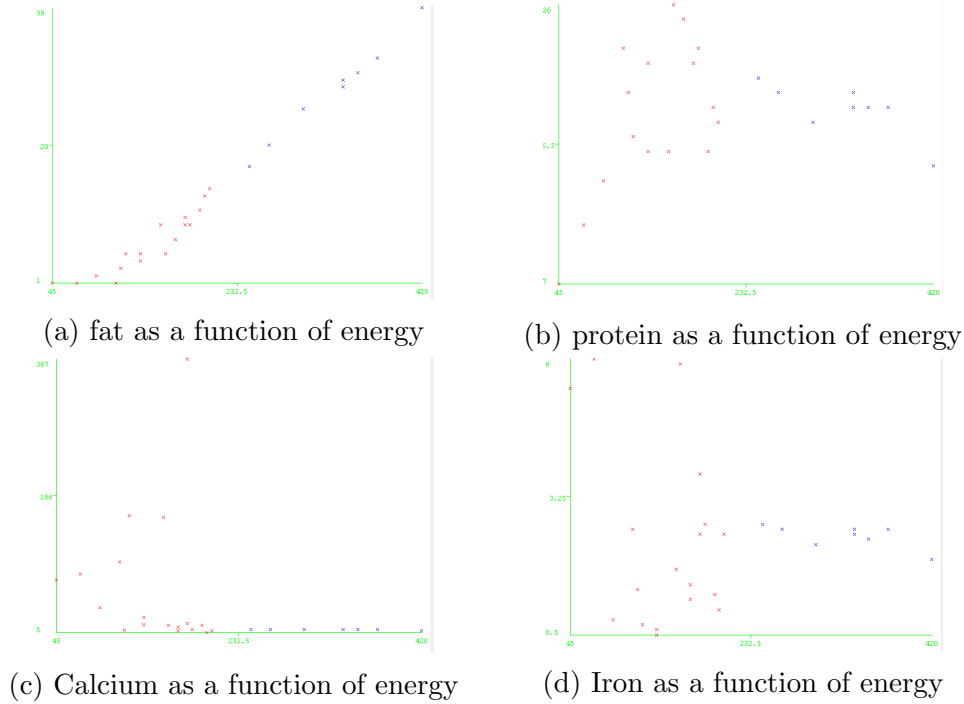


Figure 2: Cluster plot for Kmeans for 2 clusters

It is difficult to assign proper names to these clusters from the data shown in the cluster plots. However, we know that the data itself is of seafood and meats and we observed a linear relationship between the fat and energy parameters. We know that red meats contain more fat than seafood. One would push forward that we could separate cluster-wise according to this principle. Further investigation of the dataset is needed.

2.2 5 clusters with seed 10

The clusters include all the points, but some of the elements of the same clusters are very dissimilar. In Fig. 4d we can see that the elements of the green cluster are dissimilar to each other when observing iron with respect to energy. This intra-cluster dissimilarity is observed for other variables. These are bad clusters as they have high dissimilarity between elements compared to the previous solution with 2 clusters. We can conclude that one must be careful when choosing the appropriate number of clusters in order for this algorithm to work.

A special attention should be given to the type of data that is being analyzed. As noted previously, we know from the name of the data that it can be separated by red meats and seafood. There might be other ways of separating the data differently, but it is not immediately evident what 5 clusters could be made from the data.

```

=== Model information ===

Filename:      kmeans_seed10_5cl.model
Scheme:weka.clusterers.SimpleKMeans -N 5 -A "weka.core.EuclideanDistance -R first-last" -I 500 -S 10
Relation:food
Attributes:6
    Energy
    Protein
    Fat
    Calcium
    Iron
Ignored:
    Name

=== Clustering model ===

kMeans
=====

Number of iterations: 4
Within cluster sum of squared errors: 2.750432407251998
Missing values globally replaced with mean/mode

Cluster centroids:

```

Attribute	Full Data (27)	Cluster# 0 (7)	1 (8)	2 (6)	3 (1)	4 (5)
Energy	207.4074	352.8571	153.125	102.5	180	222
Protein	19	18.5714	23.25	13.5	22	18.8
Fat	13.4815	30.1429	5.75	3.8333	9	15
Calcium	43.963	8.7143	23.75	87.5	367	8.8
Iron	2.3815	2.4143	2.45	2.5333	2.5	2.02

Figure 3: Kmeans algorithm output from Weka for 5 clusters, using seed 10.

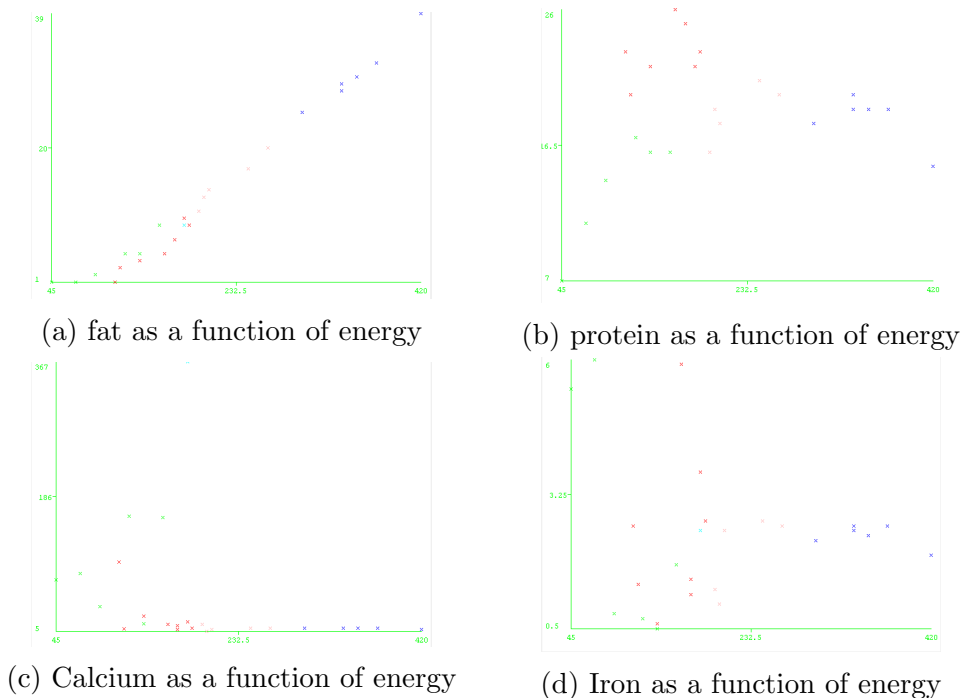


Figure 4: Cluster plot for Kmeans with 5 clusters

2.3 2 clusters with seed 1 and 123

We observe poor clustering when using seed 123, the elements within the clusters are dissimilar to one another as seen in Fig. 5c as attested by the larger sum of square error compared to the two other clustering solutions. We also notice that there are only two element in cluster 0.

When using different seed, we obtain similar, plausible results. For example, seed 1 yields similar results than those obtained for seed 10. This reinforces our belief that the clustering algorithms yields the correct result given that the other parameters are fixed (number of clusters and data). Our conclusion that the food be separated into 2 categories is reinforced, however it becomes more difficult to decide which one of these clustering solutions is the correct one. The metric we choose to separate the better clustering is the sum of square error in this case. We will choose seed 10.

The seed value controls the generation of random numbers. This in turn decides where the starting centroid is chosen to be in the algorithm. Different seed number will yield different clusters. The seed is intended to be used for reproducibility of results by different users.

```

=== Model information ===
Filename: kmeans_seed1_01.model
Scheme: weka.clusterers.SimpleKMeans -R 2 -A "weka.core.EuclideanDistance -R first-last" -I 500 -S 1
Relation: food
Attributes: {
  Energy
  Protein
  Fat
  Calcium
  Iron
Ignored: {
  Name
}

=== Clustering model ===

Ukmeans
=====

Number of iterations: 6
Within cluster sum of squared errors: 5.910434390960226
Missing values globally replaced with mean/mode

Cluster centroids:
Attribute Full Data Clusters
(27) (11) (16)
-----
Energy 207.4074 122.7273 265.626
Protein 19 14.7273 20.5625
Fat 13.4815 4.6364 19.5625
Calcium 42.563 84.2727 9.375
Iron 2.3915 2.1455 2.5625

```

(a) Kmeans using seed 1

```

=== Model information ===
Filename: kmeans_seed123_01.model
Scheme: weka.clusterers.SimpleKMeans -R 2 -A "weka.core.EuclideanDistance -R first-last" -I 500 -S 123
Relation: food
Attributes: {
  Energy
  Protein
  Fat
  Calcium
  Iron
Ignored: {
  Name
}

=== Clustering model ===

Ukmeans
=====

Number of iterations: 2
Within cluster sum of squared errors: 6.539006047715576
Missing values globally replaced with mean/mode

Cluster centroids:
Attribute Full Data Clusters
(27) (2) (25)
-----
Energy 207.4074 57.5 219.4
Protein 19 9 19.5
Fat 13.4815 1 14.48
Calcium 42.563 78 42.14
Iron 2.3915 5.7 2.114

```

(b) Kmeans using seed 123

```

Filename: kmeans_seed10_01.model
Scheme: weka.clusterers.SimpleKMeans -R 2 -A "weka.core.EuclideanDistance -R first-last" -I 500 -S 10
Relation: food
Attributes: {
  Energy
  Protein
  Fat
  Calcium
  Iron
Ignored: {
  Name
}

=== Clustering model ===

Ukmeans
=====

Number of iterations: 2
Within cluster sum of squared errors: 5.0692133929419
Missing values globally replaced with mean/mode

Cluster centroids:
Attribute Full Data Clusters
(27) (9) (18)
-----
Energy 207.4074 331.1111 145.5556
Protein 19 19 19
Fat 13.4815 27.5556 6.4444
Calcium 42.563 9.7778 61.5556
Iron 2.3915 2.4667 2.3333

```

(c) Kmeans using seed 10

Figure 5: Kmeans algorithm for 2 clusters using different seed

3 Using density based clusters

We use here density based clusters to cluster our data. This approach allows us to control how large we allow clusters to vary from their centroid by tuning the standard deviation parameter. We use the SimpleKMeans clusterer with $k = 2$ and seed 10. We will change use values of the standard deviation of $1.e^{-6}$ and 200.

We start with the density based clustering that uses the smallest standard deviation first.

```

Cluster: 0 Prior probability: 0.3448
Attribute: Energy
Normal Distribution: Mean = 331.1111 StdDev = 50.9761
Attribute: Protein
Normal Distribution: Mean = 19 StdDev = 1.63
Attribute: Fat
Normal Distribution: Mean = 27.5556 StdDev = 6.0939
Attribute: Calcium
Normal Distribution: Mean = 9.7778 StdDev = 0.6285
Attribute: Iron
Normal Distribution: Mean = 2.4667 StdDev = 0.2

Cluster: 1 Prior probability: 0.6552
Attribute: Energy
Normal Distribution: Mean = 145.5556 StdDev = 44.9245
Attribute: Protein
Normal Distribution: Mean = 19 StdDev = 4.9777
Attribute: Fat
Normal Distribution: Mean = 6.4444 StdDev = 3.8992
Attribute: Calcium
Normal Distribution: Mean = 61.5556 StdDev = 86.6962
Attribute: Iron
Normal Distribution: Mean = 2.3333 StdDev = 1.749

Time taken to build model (Full training data) : 0 seconds


=== Model and evaluation on training set ===

Clustered Instances
0 10 ( 37%)
1 17 ( 63%)

Log likelihood: -16.97893

```

(a) Density based clustering output



(b) Cluster plot of energy with respect to fat

Figure 6: Density based clustering with standard deviation of $1.e^{-6}$

We obtain similar results to the ones we previously had using the kmeans approach. Furthermore, we can observe that the clustering by visualizing fat with respect to energy in Fig. 7b is similar to that obtained previously in Fig. 2a

We now increase the standard deviation to 200 while keeping all other arguments the same.

```
Cluster: 0 Prior probability: 0.3448

Attribute: Energy
Normal Distribution. Mean = 331.1111 StdDev = 200
Attribute: Protein
Normal Distribution. Mean = 19 StdDev = 200
Attribute: Fat
Normal Distribution. Mean = 27.5556 StdDev = 200
Attribute: Calcium
Normal Distribution. Mean = 8.7778 StdDev = 200
Attribute: Iron
Normal Distribution. Mean = 2.4667 StdDev = 200

Cluster: 1 Prior probability: 0.6552

Attribute: Energy
Normal Distribution. Mean = 145.5556 StdDev = 200
Attribute: Protein
Normal Distribution. Mean = 19 StdDev = 200
Attribute: Fat
Normal Distribution. Mean = 6.4444 StdDev = 200
Attribute: Calcium
Normal Distribution. Mean = 61.5556 StdDev = 200
Attribute: Iron
Normal Distribution. Mean = 2.3389 StdDev = 200


Time taken to build model (full training data) : 0.01 seconds

=== Model and evaluation on training set ===

Clustered Instances

0      2 ( 7%)
1     25 ( 93%)

Log likelihood: -31.36038
```



(b) Cluster plot of energy with respect to fat

(a) Density based clustering output

Figure 7: Density based clustering with standard deviation of 200

We observe from the data in Fig. 7 that the when the standard deviation is large enough, we have clusters with large dissimilarity returned to us by the algorithm. As pointed out earlier, a high standard deviation parameter will result in large amounts of data that are dissimilar to the centroid to be accepted as part of the cluster. This is due to the fact that with a high enough standard deviation, the density for elements that are far from the centroid is high enough to be accepted in the first cluster that is calculated.

In conclusion, we should pay attention to the value we assign to standard deviation as a value that is too high will yield poor quality clusters with dissimilar items within them.